

Scaling Work on Odyssey

Bob Freeman, PhD
RC Facilitator
XSEDE Campus Champion

robert_freeman@harvard.edu
@DevBioInfoGuy



Scaling work on Odyssey

A few words of advice for attempting large-scale cluster work:

1. Use the appropriate partition for your job
2. Don't run large numbers of jobs out of home or lab directories; use `/scratch` or `/n/regal` instead
3. Ensure your jobs run for at least 5 - 10 minutes and keep job counts in a reasonable range (≤ 1000). Sweet spot is between 1 and 6-12 hours.
4. Test and then release incrementally: one job, several jobs, 10 jobs, more
5. Debug using interactive sessions.
6. Try to use job arrays – these are kinder on the scheduler
7. Try to use one SLURM script that accepts parameters, instead of writing out hundreds of submission scripts

